

Mahdi M.

mahdix@gmail.com
mahdix.com
linkedin.com/in/mahdix
github.com/mahdix
+44 (0)7376 503992

Profile

Hands-on technology leader with 10+ years delivering mission-critical systems for Goldman Sachs, Brevan Howard, and high-growth fintech companies. Expert in building scalable architectures using cloud-native technologies. Proven ability to lead technical transformations, from migrating legacy systems to architecting AI-enabled platforms.

Skills

- **Java, React + TypeScript, Kotlin, Python**
- **Spring, GraphQL, Thrift**
- **RESTful API, Design Patterns**
- **Kafka, Flink, Elasticsearch, PostgreSQL**
- **AWS, Docker, Kubernetes, GCP, CI/CD**
- **OOP, Functional Programming**
- Excellent analytical skills with a solid understanding of software development lifecycle, **data structures** and **algorithms**
- Extensive experience in **distributed systems** with solid understanding of **multi-threading** techniques to efficiently handle high volumes of data

Open Source (See my GitHub for more)

- **Kubernetes**: An open-source system for automating deployment, scaling, and management of containerized applications. Contributed to the “kubectf” command-line tool and Minikube (A tool that makes it easy to run Kubernetes locally).
- **Apache Cassandra**: A distributed NoSQL database with no single point of failure, based on Google Bigtable and Amazon DynamoDB. Fixed bugs in the “cqslsh” Python utility and statistics module, optimised database repair process and query tracing.
- **Solidinbox**: A personal side-project for lead generation and marketing on Twitter (Java + MySQL + React + Typescript).

Professional Experience

Co-founder and Chief Technology Officer (Dec '24 - present) **FactFolio** (London, United Kingdom)

- Worked on an AI-enabled portfolio monitoring SaaS aimed at infrastructure investors that handles data collection, quality checks and reporting automation (MS Word and MS PowerPoint export) while providing early warning alerts for portfolio underperformance.
- Used the OpenAI API to enable users to chat with their data and ask questions.
- Technology Stack: React + Typescript / Java + PostgreSQL / AWS / Caddy

Systematic Engineer (Dec '22 - Dec '24) **Coremont.com/Brevan Howard** (London, United Kingdom)

- Worked on a high-frequency, scalable and distributed post-trade processing system to accept incoming algo trades (across different markets like Equities, Futures, Swaps and Bonds), process and aggregate them and book them into our order management system.
- Worked on cash flow calculation and booking for Equities and Equity Swaps and store into Murex.
- The system consisted of multiple Java-based microservices (AWS EKS), talking to each other via Kafka (AWS MKS), saving state in MongoDB or PostgreSQL and communicating with external systems via S3/Kafka/ActiveMQ + “Apache Flink” based recon pipeline.

Staff Engineer (April '22 - Dec '22) **Blockchain.com** (London, United Kingdom)

- As a hands-on member of the Core Platform team, worked on infrastructure-related projects (Ledger and SSO) to develop technical specification documents, design a solution and implement it.

Senior Software Engineer (Mar '21 - April '22) **Credit Karma** (London, United Kingdom)

- Part of the International Product team working on both UK and Canada websites as a full-stack engineer (React, Typescript, Scala, GraphQL, MySQL, Apache Thrift, K8s, GCP).

Vice President (Jan '18 - Mar'21) **Goldman Sachs** (London, United Kingdom)

- Part of the data distribution team (Investment Management division), worked on different data aggregation and processing systems.
- Complex event processing to import, extract, cleanse, enrich, aggregate, link, persist and provide data coming from different sources at different times. Exposed through Web APIs to be consumed by other systems.

Senior Software Engineer (Mar '17 - Jan '18) **Skimlinks** (London, United Kingdom)

- Worked as a member of the data-platform team with an Apache Hadoop cluster to implement big-data processing pipelines.
- Implemented Apache Oozie workflows to cleanse, aggregate and process 500 million records per hour and store them in Hive.
- A RESTful API (Python Flask + uWSGI + nginx) for Elasticsearch on a K8s cluster hosted on Google Cloud for reporting systems.

Software Developer (Oct '14 – Mar '17) [Binary.com](#) (Cyberjaya, Malaysia)

- Migrated and refactored the legacy Apache CouchDB-based financial data system (for Interest Rate, Volatility Surface, Spot, ...) to a new storage model (Redis with distributed replication across all servers + replicated PostgreSQL). Achieved higher performance and lowered maintenance costs, without interrupting the 1.2M daily transactions (providing 24/7 service).
- Developed code for real-time financial data processing from different providers ([Bloomberg](#), [ICE](#), [SIX iD](#) and [SuperDerivatives](#)).